

# Lecture 4

---

3B1B Optimization

Michaelmas 2015

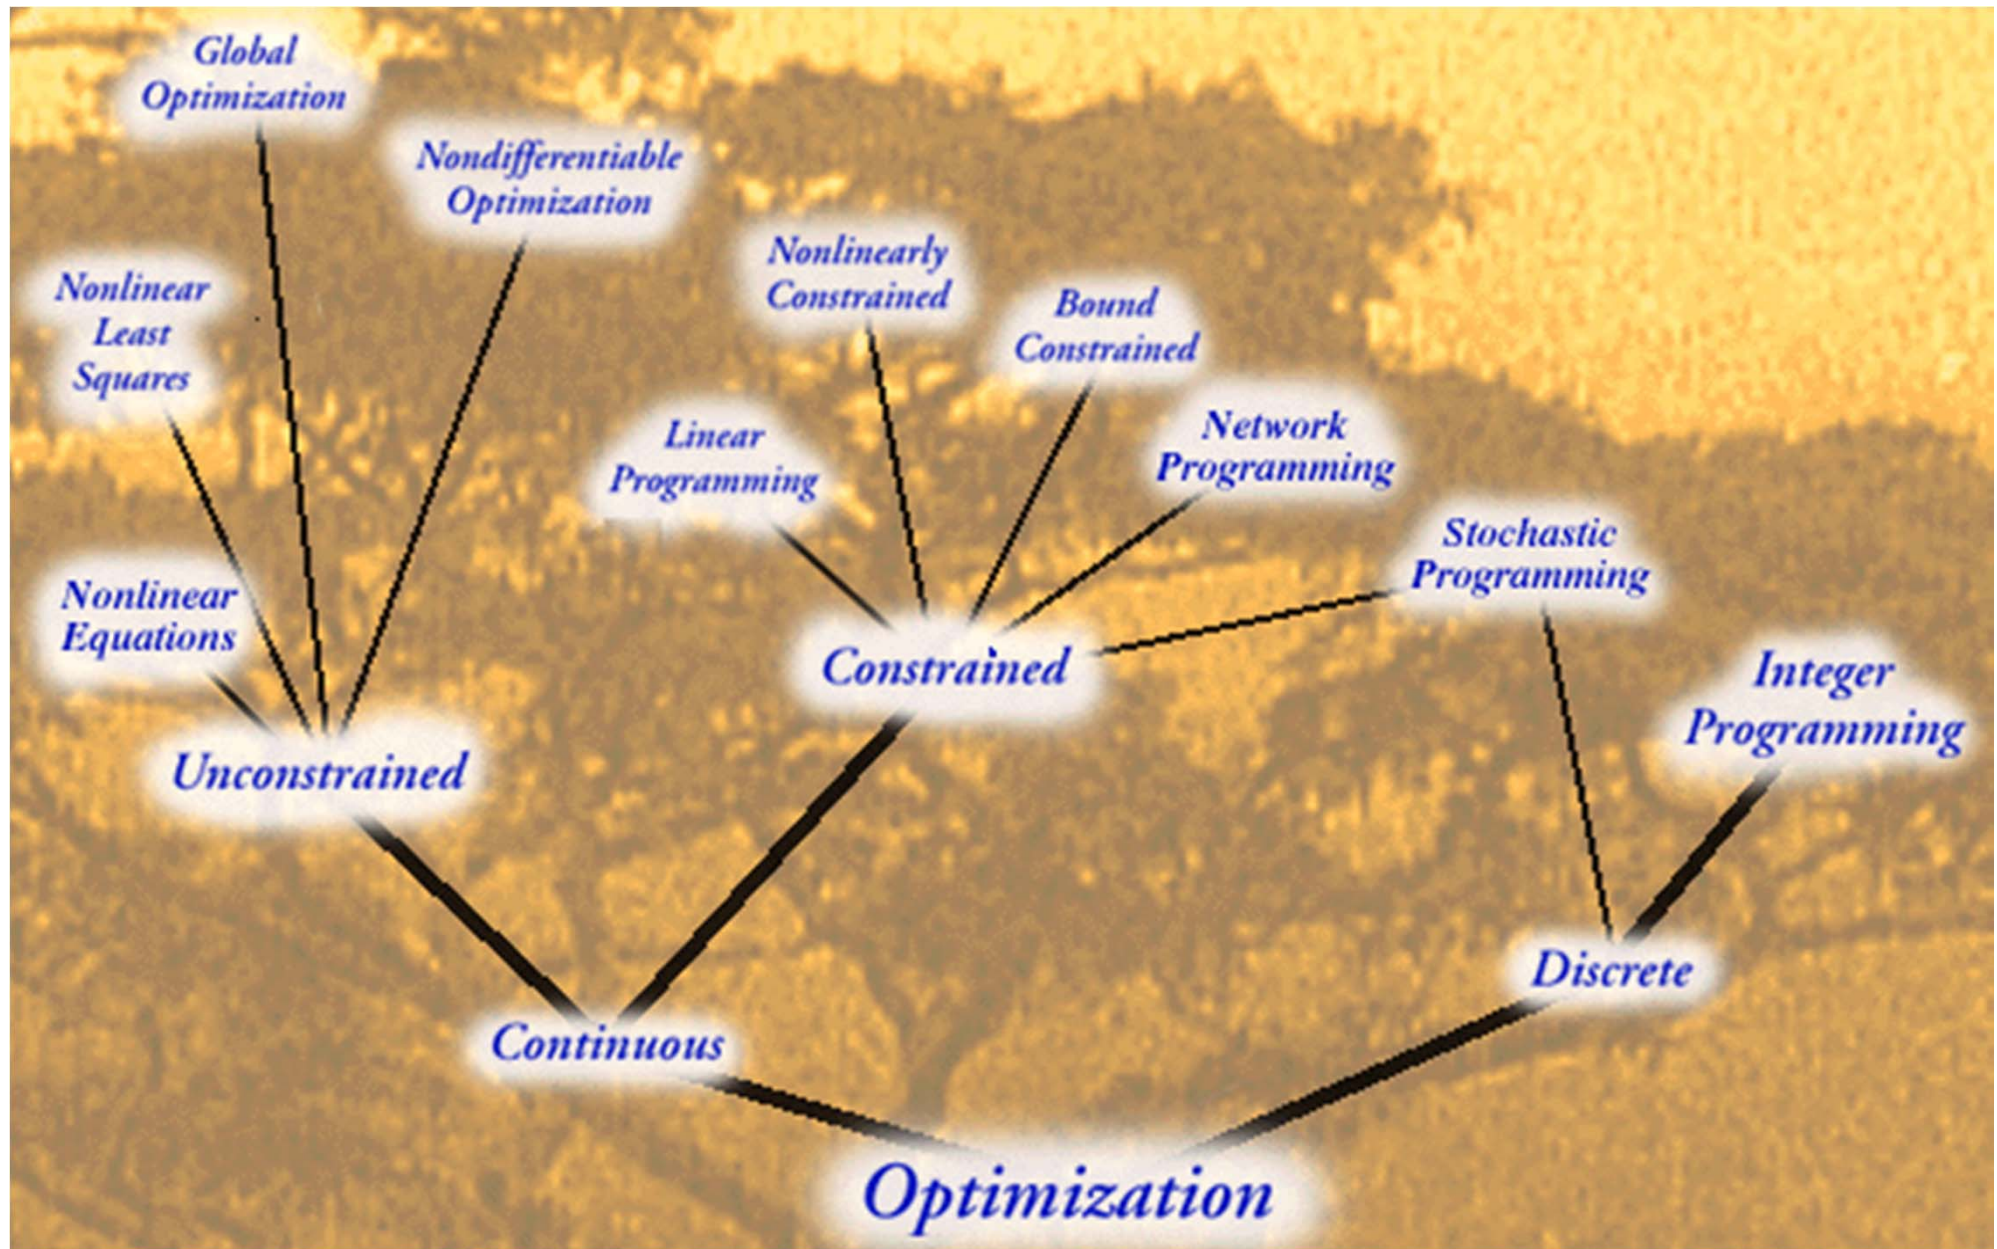
A. Zisserman

---

- Convexity
- Robust cost functions
- Optimizing non-convex functions
  - grid search
  - branch and bound
  - multiple coverings
  - simulated annealing

# The Optimization Tree

---




# Unconstrained optimization

---

function of one  
variable


$$\min_x f(x)$$




- down-hill search (gradient descent) algorithms can find local minima
- which of the minima is found depends on the starting point
- such minima often occur in real applications

# Cost functions in 2D

---



Global maximum



Multiple optima

# How can you tell if an optimization has a single optimum?

---

The answer is: see if the optimization problem is **convex**.  
If it is, then a local optimum is the global optimum.


First, we need to introduce

- Convex Sets, and
- Convex Functions


Note – sketch introduction only

# Convex set

A set  $D \subset \mathbb{R}^n$  is **convex** if the line joining points  $x_0$  and  $x_1$  lies inside  $D$ .



Intersection of convex sets is convex.






# Convex functions

$D$  – a domain in  $\mathbb{R}^n$ .


A **convex function**  $f : D \rightarrow \mathbb{R}$  is one that satisfies, for any  $\mathbf{x}_0$  and  $\mathbf{x}_1$  in  $D$ :

$$f((1 - \alpha)\mathbf{x}_0 + \alpha\mathbf{x}_1) \leq (1 - \alpha)f(\mathbf{x}_0) + \alpha f(\mathbf{x}_1) .$$



Line joining  $(\mathbf{x}_0, f(\mathbf{x}_0))$   
and  $(\mathbf{x}_1, f(\mathbf{x}_1))$  lies  
above the function graph.



# Convex function examples



convex



Not convex

A non-negative sum of convex functions is convex



# Convex Optimization Problem

---

Minimize:


- a convex function
- over a convex set

Then locally optimal points are globally optimal

Also, such problems can be solved both in theory and practice

# Why do we need the domain to be convex?

---



# Examples of convex optimization problems

---

1. Linear programming

2. Least squares

$$f(\mathbf{x}) = (\mathbf{Ax} - \mathbf{b})^2, \text{ for any } \mathbf{A}$$

3. Quadratic functions

$$f(\mathbf{x}) = \mathbf{x}^\top \mathbf{P} \mathbf{x} + \mathbf{q}^\top \mathbf{x} + r, \text{ provided that } \mathbf{P} \text{ is positive definite}$$

Many more useful examples, see Boyd & Vandenberghe

## First-order condition


$f$  is **differentiable** if  $\text{dom } f$  is open and the gradient

$$\nabla f(x) = \left( \frac{\partial f(x)}{\partial x_1}, \frac{\partial f(x)}{\partial x_2}, \dots, \frac{\partial f(x)}{\partial x_n} \right)$$

exists at each  $x \in \text{dom } f$

**1st-order condition:** differentiable  $f$  with convex domain is convex iff

$$f(y) \geq f(x) + \nabla f(x)^T (y - x) \quad \text{for all } x, y \in \text{dom } f$$



first-order approximation of  $f$  is global underestimator

# Second order condition

---

The Hessian of a function  $f(x_1, x_2, \dots, x_n)$  is the matrix of partial derivatives

$$H = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1 \partial x_1} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n \partial x_n} \end{bmatrix}$$

Diagonalize the Hessian by an orthogonal change of coordinates.

Diagonals are the **eigenvalues**.


If the eigenvalues are all **positive**, then the Hessian is **positive definite**, and  $f$  is **convex**.

# Strictly convex


---

A function  $f(x)$  is **strictly convex** if

$$f((1 - \alpha)\mathbf{x}_0 + \alpha\mathbf{x}_1) < (1 - \alpha)f(\mathbf{x}_0) + \alpha f(\mathbf{x}_1) .$$



strictly convex  
one global optimum



**Not strictly convex**  
multiple local optima  
(but all are global)

# Robust Cost Functions

- In formulating an optimization problem there is often some room for design and choice
- The cost function can be chosen to be:
  - convex
  - robust to noise (outliers) in the data/measurements




# Motivation

---

## Fitting a 2D line to a set of 2D points


- Suppose you fit a straight line to data containing **outliers** – points that are not properly modelled by the assumed measurement noise distribution
- The usual method of least squares estimation is hopelessly corrupted



# Least squares: Robustness to noise

---


Least squares fit to the red points:



# Least squares: Robustness to noise



---

Least squares fit with an outlier:




Problem: squared error heavily penalizes outliers

Consider minimizing the cost function  $f(x) = \sum_i (x - a_i)^2$




- the data  $\{a_i\}$  may be thought of as repeated measurements of a fixed value (at 0), subject to Gaussian noise and some outliers
- it has 10% of outliers biased towards the right of the true value
- the minimum of  $f(x)$  does not correspond to the true value


Examine the behaviour of various cost functions  $f(x) = \sum_i C(|x - a_i|)$




quadratic



truncated quadratic



$L_1$




huber

# Quadratic cost function

---

- squared error – the usual default cost function
- arises in Maximum Likelihood Estimation for Gaussian noise
- convex




$$C(\delta) = \delta^2$$

# Truncated Quadratic cost function

---

- for inliers behaves as a quadratic
- truncated so that outliers only incur a fixed cost
- non-convex



$$C(\delta) = \min(\delta^2, \alpha)$$


$$= \begin{cases} \delta^2 & \text{if } |\delta| < \sqrt{\alpha} \\ \alpha & \text{otherwise.} \end{cases}$$



# $L_1$ cost function

---

- absolute error
- called 'total variation'
- convex
- non-differentiable at origin
- finds the **median** of  $\{a_i\}$




$$C(\delta) = |\delta|$$

# Huber cost function


---

- hybrid between quadratic and  $L_1$
- continuous first derivative
- for small values is quadratic
- for larger values becomes linear
- thus has the outlier stability of  $L_1$
- convex




$$C(\delta) = \begin{cases} \delta^2 & \text{if } |\delta| < \beta \\ 2\beta|\delta| - \beta^2 & \text{otherwise.} \end{cases}$$


## Example 1: measurements with outliers




$$f(x) = \sum_i C(|x - a_i|)$$




quadratic



truncated quadratic




huber




zoom

## Example 2: bimodal measurements


- 70% in principal mode
- 30% in outlier mode




$$f(x) = \sum_i C(|x - a_i|)$$



quadratic



truncated quadratic



huber

# Summary


---

- Squared cost function very susceptible to outliers
- Truncated quadratic has a stable minimum, but is non-convex and also has other local minima. Also basin of attraction of global minimum limited
- Huber has stable minimum and is convex



# Application 1: Robust line fitting

---



(See also RANSAC algorithm)

Boyd & Vandenberghe



# Application 2: Signal restoration


---

Measurements  $z_i$  are original signal  $x_i$  corrupted with additive noise

$$z_i = x_i + w_i$$

where

$$w_i \sim N(0, \sigma^2)$$





Compare:

$$f(\mathbf{x}) = \sum_i (z_i - x_i)^2 + \lambda (x_i - x_{i-1})^2 \quad \text{Quadratic smoothing}$$

$$f(\mathbf{x}) = \sum_i (z_i - x_i)^2 + \lambda |x_i - x_{i-1}| \quad \text{Total variation}$$

## Quadratic smoothing

$$f(\mathbf{x}) = \sum_i (z_i - x_i)^2 + \lambda (x_i - x_{i-1})^2$$








Quadratic smoothing smooths out noise **and** steps in the signal

Boyd & Vandenberghe

## Total variation

$$f(\mathbf{x}) = \sum_i (z_i - x_i)^2 + \lambda |x_i - x_{i-1}|$$



Total variation smoothing preserves steps in the signal


Boyd & Vandenberghe

# Optimizing non-convex functions

---

function of one  
variable

$$\min_x f(x)$$




Sketch four methods:

1. grid search: uniform grid space covering
2. branch and bound
3. multiple coverings: Newton like methods within regions
4. simulated annealing: stochastic optimization

# Branch and bound

---



Key idea:



- Split region into sub-regions and compute bounds
- Consider two regions A and C
- If lower bound of A is greater than upper bound of C then A can be discarded
- divide (branch) regions and repeat

# Multiple coverings

---

Key idea is to cover the parameter space with overlapping regions to deal with local optima, and then take advantage of efficient continuous optimization for each region.


## Example from Matlab Global Optimization toolbox



# Multiple coverings ctd

---


- Use multiple starting points
- Continuous optimization method for each
- Record optimum for each starting point
- Sort values to find global optimum





# Simulated Annealing

---



- The algorithm has a mechanism to jump out of local minima
- It is a **stochastic search** method, i.e. it uses randomness in the search

# Simulated annealing algorithm

---

- At each iteration propose a move in the parameter space
- If the move decreases the cost, then accept it
- If the move increases the cost by  $\Delta E$ , then
  - accept it with a probability  $\propto \exp(-\Delta E/T)$ ,
  - Otherwise, don't move

Note probability depends on **temperature T**

- Decrease the temperature according to a **schedule** so that at the start cost increases are likely to be accepted, and at the end they are not

# Boltzmann distribution and the cooling schedule

---

- start with  $T$  high, then  $\exp(-\Delta E/T)$  is approx. 1, and all moves are accepted


- many cooling schedules are possible, but the simplest is

$$T_{k+1} = \alpha T_k, \quad 0 < \alpha < 1$$

where  $k$  is the iteration number

- The algorithm can be very slow to converge ...

Boltzmann distribution  $\exp(-\Delta E/T)$



# Simulated annealing

---


The name and inspiration come from annealing in metallurgy, a technique involving heating and controlled cooling of a material to increase the size of its crystals and reduce their defects.

The heat causes the atoms to become unstuck from their initial positions (a local minimum of the internal energy) and wander randomly through states of higher energy; the slow cooling gives them more chances of finding configurations with lower internal energy than the initial one.

Algorithms due to: Kirkpatrick *et al.* 1982; Metropolis *et al.* 1953.


# Example: Convergence of simulated annealing

---




# Steepest descent on a graph

---




# Random Search on a graph

---



# Simulated Annealing on a graph

---





# Simulated annealing for the Rosenbrock function

---




Figure 18: Minimization of the two-dimensional Rosenbrock function by simulated annealing--search pattern.




Figure 19: Minimization of the two-dimensional Rosenbrock function by simulated annealing objective reduction.

# There is more ...

---

There are many other classes of optimization problem, and also many efficient optimization algorithms developed for problems with special structure. Examples include:

- Combinatorial and discrete optimization
- Dynamic programming
- Max-flow/Min-cut graph cuts
- ...

See the links on the web page

<http://www.robots.ox.ac.uk/~az/lectures/b1/index.html>

and come to the C Optimization lectures next year